
A Pythonic Job Search

— —

Sam Raykhman

Job Searching CAN be Stressful

Objective: To establish a model that classifies jobs that have knowledge of Python as a requirement

Purpose: Help out myself and my peers to find a job for our qualifications



The Process

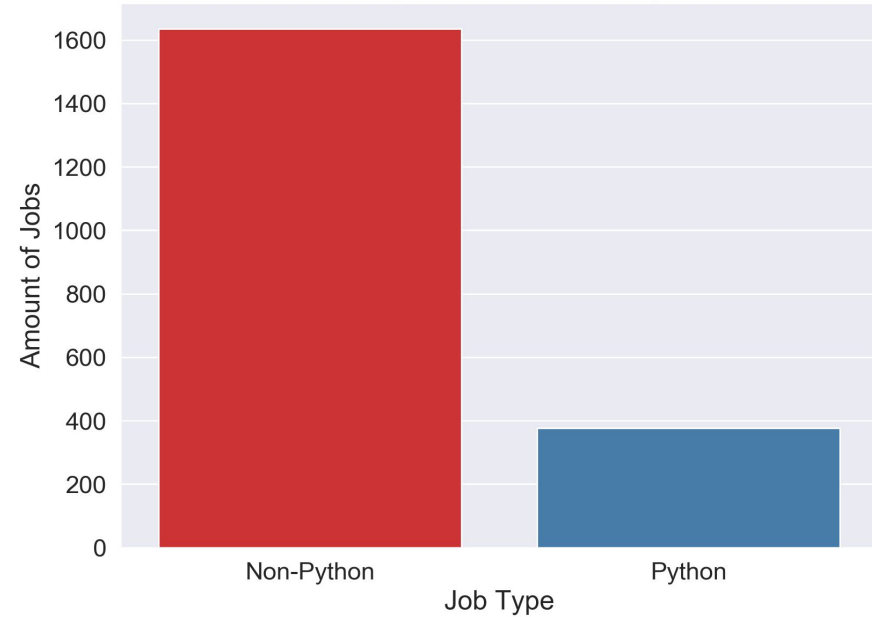
Data Collection	EDA	NLP Feature Engineering	Modeling	Evaluation
Web scraped my data from <u>stackoverflow</u>	Cursory examination of data	Tokenized, removed stop words and performed TF-IDF	Multinomial Naive Bayes, Random Forest, XGBoost	Do I need python or not?

Diving into the Data

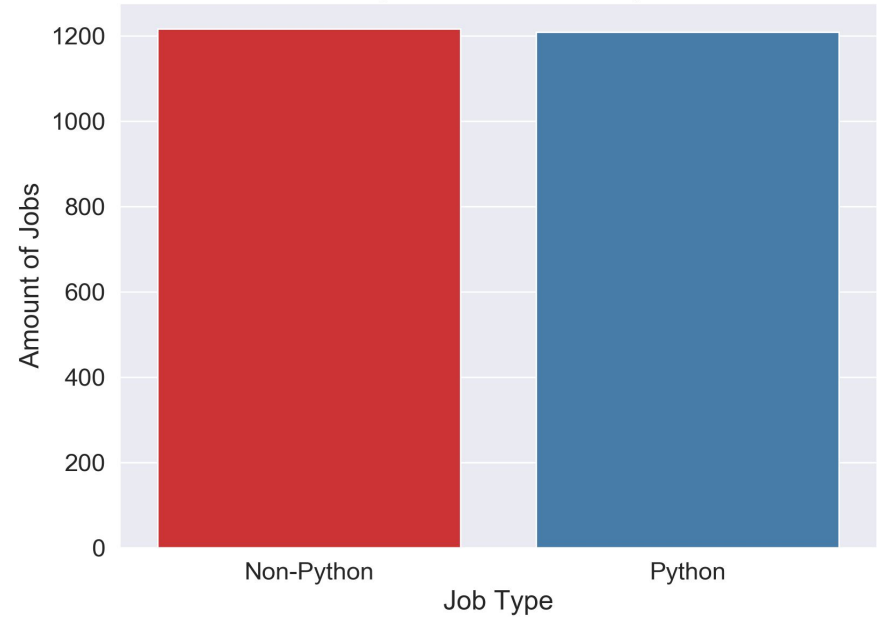
- I web scraped my data from stackoverflow's job posting portal
 - At the time of my scraping there were 2893 jobs in total
 - After removing rows with empty descriptions or no list of programming languages, I was left with about 2100 rows.
- I created 5 features
 - Job Title
 - Description
 - Languages
 - Overview
 - URL
- Created target variable out of Languages column
 - If python was listed as a language required or not

Class Imbalance

Unbalanced: Python Jobs vs. Non-Python Jobs




Balanced: Python Jobs vs. Non-Python Jobs



Common Languages

Count of Most Common Programming Languages




Word Clouds

20 most common words for jobs
without Python as a requirement



20 most common words for jobs with
Python as a requirement



Model Evaluation

Dummy

Accuracy: 0.8330

F1 Score: 0.0



- Strategy = most frequent

Naive Bayes

Accuracy: 0.8489

F1 Score: 0.4933



- Alpha = 0.0001

Random Forest

Accuracy: 0.8330

F1 Score: 0.0



- Criterion = gini
- Max depth = 1
- Max features = 1
- Max leaf nodes = 2
- # of estimators = 20

XGBoost

Accuracy: 0.8887

F1 Score: 0.6627



- Learning rate = 0.1
- Max depth = 3
- Min child weight = 1
- # of estimators = 200

ADASYN XGBoost


Accuracy: 0.8429

F1 Score: 0.6291




- Learning rate = 0.001
- Max depth = 5
- Min child weight = 0.01
- # of estimators = 100

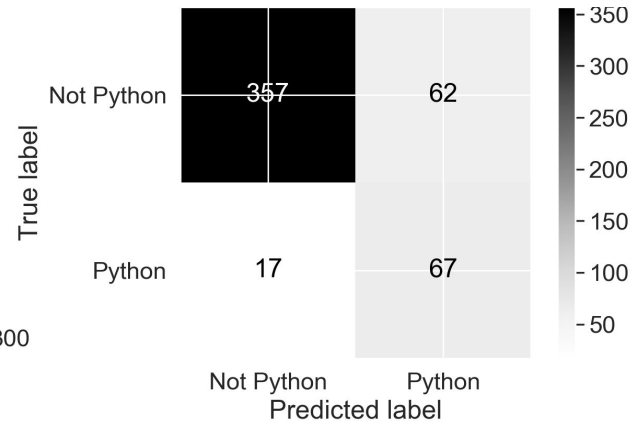
Confusion Matrices



Naive Bayes

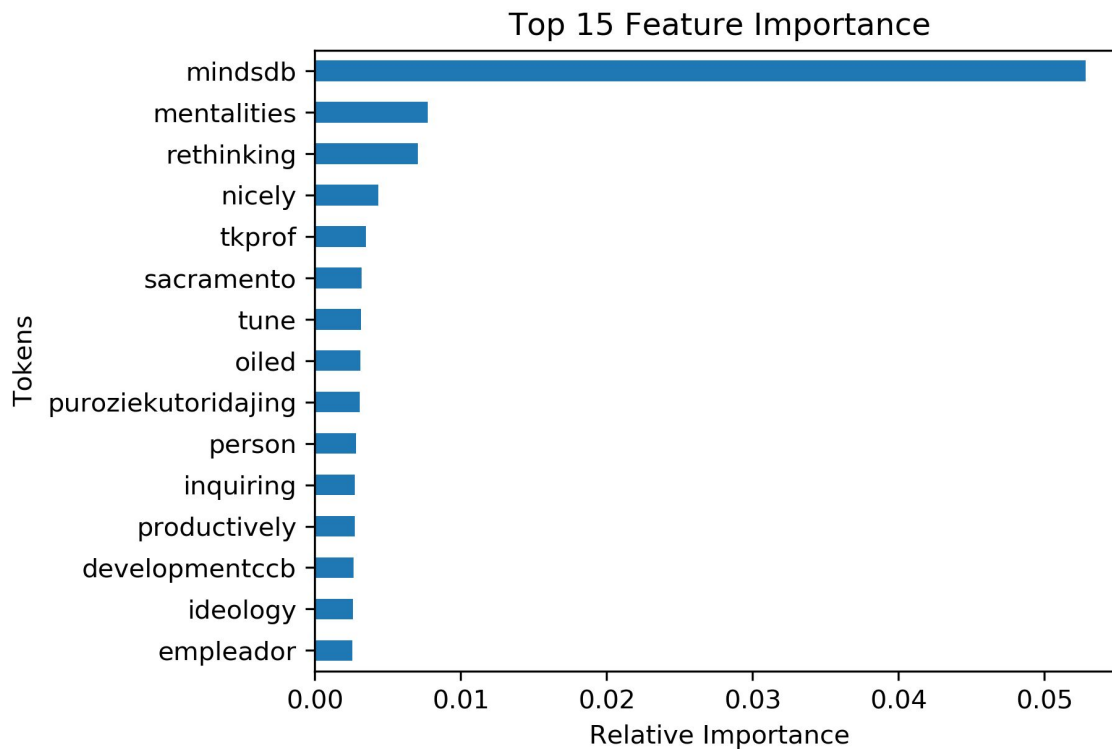


Unbalanced XGBoost



XGBoost with ADASYN

RF Feature Importance



Conclusion

In the real world, it can take hours to merely find a job that you qualify for and interests you. Let alone, the days and weeks of applying and going through the interview processes.

Using my unbalanced XGBoost model, I can easily filter the jobs to my (or other people's) qualifications, reducing the time and effort spent on the hunt; without hindering the results by generating false data points through SMOTE or ADASYN.

Future steps

- Incorporate N-grams, to check importance of phrases and not just words
- Use a Support Vector Machine Model, to attempt to increase the performance of my final model
- Create recommendation system
 - That would take qualifications, location, and some consumer preferences to return jobs in your area that match your requests

Thank you

Any Questions?